

```

from google.colab import drive
drive.mount('/content/drive')

import os

# TODO: Fill in the Google Drive path where you uploaded the lab materials
# Example: GOOGLE_DRIVE_PATH_AFTER_MYDRIVE = 'Colab Notebooks'

GOOGLE_DRIVE_PATH_AFTER_MYDRIVE = 'Colab Notebooks/Computer
Vision/CW_Folder_PG/Code'
GOOGLE_DRIVE_PATH = os.path.join('drive', 'My Drive',
GOOGLE_DRIVE_PATH_AFTER_MYDRIVE)
print(os.listdir(GOOGLE_DRIVE_PATH))

Preprocessing -

import os
import numpy as np
from PIL import Image

def preprocess_dataset(images_path, labels_path, target_size=(128, 128),
convert_gray=True):
    images_list = [] # List to hold preprocessed images
    labels_list = [] # List to hold corresponding labels

    # Process each file in the directory
    for filename in os.listdir(images_path):
        if filename.lower().endswith(('.jpg', '.jpeg', '.png')):
            img_path = os.path.join(images_path, filename)
            img = Image.open(img_path).resize(target_size) # Open and resize
image

            # Convert to grayscale if specified
            if convert_gray and img.mode != 'L':
                img = img.convert('L')

            img_array = np.array(img) / 255.0 # Normalize pixel values to [0,
1]

            # Add channel dimension to grayscale images
            if convert_gray or img.mode == 'L':
                img_array = img_array[..., np.newaxis]

            images_list.append(img_array) # Add processed image to list

            # Load and append label
            label_filename = os.path.splitext(filename)[0] + '.txt'
            label_path = os.path.join(labels_path, label_filename)
            with open(label_path, 'r') as file:
                label = int(file.read().strip())
            labels_list.append(label)

    return np.array(images_list), np.array(labels_list) # Return arrays of
images and labels

```

```

# Paths to your datasets
train_images_path = '/content/drive/MyDrive/Colab Notebooks/Computer
Vision/CW_Folder_PG/CW_Dataset/train/images'
train_labels_path = '/content/drive/MyDrive/Colab Notebooks/Computer
Vision/CW_Folder_PG/CW_Dataset/train/labels'
x_train, y_train = preprocess_dataset(train_images_path, train_labels_path)

test_images_path = '/content/drive/MyDrive/Colab Notebooks/Computer
Vision/CW_Folder_PG/CW_Dataset/test/images'
test_labels_path = '/content/drive/MyDrive/Colab Notebooks/Computer
Vision/CW_Folder_PG/CW_Dataset/test/labels'
x_test, y_test = preprocess_dataset(test_images_path, test_labels_path)

print("Preprocessing completed for both training and test datasets.")

import matplotlib.pyplot as plt

# Display the first 10 preprocessed images in a 2x5 grid
fig, axes = plt.subplots(2, 5, figsize=(15, 8))

for i in range(10):
    row = i // 5
    col = i % 5
    # If the image is 3D (has a channel dimension), but is still grayscale,
    # we convert it to 2D for displaying.
    image = x_train[i].squeeze() # This removes the channel dim if it's 1
    axes[row, col].imshow(image, cmap='gray') # Force grayscale colormap
    axes[row, col].set_title(f"Label: {y_train[i]}")
    axes[row, col].axis('off')

plt.tight_layout()
plt.show()

import numpy as np
import matplotlib.pyplot as plt

# Assuming 'y_train' contains class labels
unique_classes, class_counts = np.unique(y_train, return_counts=True)

# Plot class distribution
plt.bar(unique_classes, class_counts)
plt.xlabel('Class')
plt.ylabel('Number of Instances')
plt.title('Class Distribution')
plt.xticks(unique_classes)
plt.show()

# Calculate class proportions
total_samples = len(y_train)
class_proportions = class_counts / total_samples

# Print class proportions
for cls, prop in zip(unique_classes, class_proportions):

```

```
print(f"Class {cls}: Proportion = {prop:.2f}")
```

HOG_SVM -

```
# Import necessary libraries
#Reference Lab 5
from skimage.feature import hog # Import HOG feature extraction
from sklearn import svm # Import SVM classifier
from sklearn.metrics import classification_report, accuracy_score,
confusion_matrix # Import evaluation metrics
from sklearn.model_selection import train_test_split, GridSearchCV # Import
data splitting and grid search
from tensorflow.keras.preprocessing.image import ImageDataGenerator # Import
image data augmentation
import joblib # For saving the model
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler # Import feature scaling

# Function to extract HOG features from images
def extract_hog_features(images):
    hog_features = []
    for image in images:
        if image.ndim > 2:
            image = image[:, :, 0] # Convert to grayscale if not already
        # Extract HOG features from the image
        fd = hog(image, orientations=8, pixels_per_cell=(16, 16),
            cells_per_block=(1, 1), block_norm='L2', visualize=False)
        hog_features.append(fd)
    return np.array(hog_features)

# Function to augment data for imbalanced classes
def augment_data(x_data, y_data, classes_to_augment=[0, 2], augment_factor=5):
    # Setup the data augmentation configuration
    data_gen = ImageDataGenerator(
        rotation_range=10, # Rotate the image within a range of 10
        degrees
        width_shift_range=0.1, # Horizontally shift the image by 10%
        height_shift_range=0.1, # Vertically shift the image by 10%
        shear_range=0.1, # Shear the image by 10%
        zoom_range=0.1, # Zoom in/out by 10%
        horizontal_flip=True, # Allow horizontal flipping of the image
        fill_mode='nearest' # Fill in new pixels with the nearest filled
        value
    )
    x_augmented = [] # List to store augmented images
    y_augmented = [] # List to store labels of augmented images
    for class_label in classes_to_augment:
        # Filter out images of the specified classes
        x_class = x_data[y_data == class_label]
        # Augment images
```

```

        for _ in range(augment_factor):
            augmented_images = data_gen.flow(x_class, batch_size=len(x_class),
shuffle=False)
            for aug_img in augmented_images:
                x_augmented.extend(aug_img)
                y_augmented.extend([class_label] * len(aug_img))
                break # we only need one set per batch
    # Combine original and augmented data
    x_new = np.concatenate((x_data, np.array(x_augmented)))
    y_new = np.concatenate((y_data, np.array(y_augmented)))
    return x_new, y_new

# Assuming x_train and y_train are loaded and preprocessed

# Augment data for specified classes and split into training and validation sets
x_train, y_train = augment_data(x_train, y_train, classes_to_augment=[0, 2],
augment_factor=5)
x_train_hog = extract_hog_features(x_train) # Extract HOG features for training
data
x_test_hog = extract_hog_features(x_test)    # Extract HOG features for test
data

# Split data into training and validation sets
x_train_hog, x_val_hog, y_train, y_val = train_test_split(x_train_hog, y_train,
test_size=0.3, random_state=42)

# Initialize and configure the scaler
scaler = StandardScaler()
x_train_hog_scaled = scaler.fit_transform(x_train_hog) # Fit and transform the
training data
x_val_hog_scaled = scaler.transform(x_val_hog)          # Transform the
validation data based on the fitted scaler

# Save the scaler for later use
joblib.dump(scaler, 'hog_scaler.pkl')

# Define the parameter grid for grid search
param_grid = {
    'C': [0.1, 1, 10, 100],
    'gamma': ['scale', 'auto'],
    'kernel': ['linear', 'rbf']
}

# Initialize and configure the SVM classifier
svm_clf = svm.SVC(class_weight='balanced') # Initialize SVM with balanced class
weights
grid_search = GridSearchCV(svm_clf, param_grid, refit=True, verbose=3, cv=5) #
Perform grid search
grid_search.fit(x_train_hog_scaled, y_train) # Note the use of scaled HOG
features here

print("Best parameters found:", grid_search.best_params_) # Print the best
parameters found
best_clf = grid_search.best_estimator_ # Get the best SVM classifier

```

```

# Predict and evaluate as before using scaled HOG features
y_train_pred = best_clf.predict(x_train_hog_scaled)
print("Training Accuracy:", accuracy_score(y_train, y_train_pred))

y_pred_val = best_clf.predict(x_val_hog_scaled)
val_accuracy = accuracy_score(y_val, y_pred_val)
print("Validation Accuracy:", val_accuracy)

conf_matrix = confusion_matrix(y_val, y_pred_val)
print("Confusion Matrix:\n", conf_matrix)

class_report = classification_report(y_val, y_pred_val)
print("Classification Report:\n", class_report)

# Visualize the confusion matrix using seaborn's heatmap
plt.figure(figsize=(6, 6))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='viridis',
xticklabels=np.unique(y_val), yticklabels=np.unique(y_val))
plt.title('Confusion Matrix for the Validation Set')
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.show()

# Save the model
joblib.dump(best_clf, 'best_hog_svm_model.pkl')

from sklearn.preprocessing import StandardScaler # Import StandardScaler for
feature scaling
from sklearn.metrics import accuracy_score, confusion_matrix,
classification_report # Import evaluation metrics
import joblib # Import joblib for loading the model and scaler

# Load the trained model and scaler
model = joblib.load('best_hog_svm_model.pkl') # Load the trained SVM model
scaler = joblib.load('hog_scaler.pkl') # Load the scaler used for
feature scaling during training

# Transform the test HOG features using the loaded scaler
x_test_scaled = scaler.transform(x_test_hog) # Scale the test features

# Predict on the test set using the trained model
y_test_pred = model.predict(x_test_scaled) # Predict the labels for the test
set

# Calculate and print the test accuracy
test_accuracy = accuracy_score(y_test, y_test_pred) # Compute the accuracy of
the predictions
print("Test Accuracy:", test_accuracy) # Print the test accuracy

# Generate and display the confusion matrix and classification report
conf_matrix = confusion_matrix(y_test, y_test_pred) # Compute the confusion
matrix
print("Confusion Matrix:\n", conf_matrix) # Print the confusion matrix

```

```

class_report = classification_report(y_test, y_test_pred) # Generate the
classification report
print("Classification Report:\n", class_report) # Print the classification
report

```

```

# Visualize the confusion matrix using seaborn's heatmap
import matplotlib.pyplot as plt
import seaborn as sns
plt.figure(figsize=(6, 6))
sns.heatmap(conf_matrix, annot=True, fmt='g', cmap='viridis',
xticklabels=np.unique(y_test), yticklabels=np.unique(y_test))
plt.xlabel('Predicted labels')
plt.ylabel('True labels')
plt.title('Confusion Matrix for Test Set')
plt.show() # Display the heatmap of the confusion matrix

```

SIFT_SVM-

```

# Import necessary libraries
import cv2
import numpy as np
from skimage import img_as_ubyte
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from sklearn.cluster import MiniBatchKMeans
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, confusion_matrix,
classification_report
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt
import seaborn as sns
import joblib

# Function to augment data for imbalanced classes
def augment_data(x_data, y_data, classes_to_augment=[0, 2], augment_factor=5):
    # Setup the data augmentation configuration
    data_gen = ImageDataGenerator(
        rotation_range=10, # Rotate the image within a range of 10 degrees
        width_shift_range=0.1, # Horizontally shift the image by 10%
        height_shift_range=0.1, # Vertically shift the image by 10%
        shear_range=0.1, # Shear the image by 10%
        zoom_range=0.1, # Zoom in/out by 10%
        horizontal_flip=True, # Allow horizontal flipping of the image
        fill_mode='nearest' # Fill in new pixels with the nearest filled value
    )
    x_augmented = [] # List to store augmented images
    y_augmented = [] # List to store labels of augmented images
    for class_label in classes_to_augment:
        # Filter out images of the specified classes
        x_class = x_data[y_data == class_label]
        x_class = x_class.reshape((x_class.shape[0],) + x_class.shape[1:])
        # Generate augmented images
        aug_iter = data_gen.flow(x_class, batch_size=1)

```

```

        for _ in range(x_class.shape[0] * augment_factor):
            x_augmented.append(next(aug_iter)[0])
            y_augmented.append(class_label)
# Combine original and augmented data
x_new = np.concatenate((x_data, np.array(x_augmented)))
y_new = np.concatenate((y_data, np.array(y_augmented)))
return x_new, y_new

# Augment data for specified classes and split into training and validation sets
x_train, y_train = augment_data(x_train, y_train, classes_to_augment=[0, 2],
                                augment_factor=5)
x_train, x_val, y_train, y_val = train_test_split(x_train, y_train,
                                                    test_size=0.2, random_state=42)

# Initialize the SIFT feature detector
sift = cv2.SIFT_create()
des_list = [] # List to store descriptors
y_train_list = [] # List to store labels corresponding to the descriptors

# Setup figure to visualize keypoints detected by SIFT
fig, ax = plt.subplots(1, 4, figsize=(10, 8), sharey=True)
for i in range(len(x_train)):
    img = img_as_ubyte(x_train[i]) # Convert the image for SIFT processing
    kp, des = sift.detectAndCompute(img, None) # Detect keypoints and compute
descriptors
    if i < 4: # Display keypoints on the first 4 images
        img_with_SIFT = cv2.drawKeypoints(img, kp, img)
        ax[i].imshow(img_with_SIFT, cmap='gray')
        ax[i].set_axis_off()
    if des is not None:
        des_list.append(des) # Store descriptors
        y_train_list.append(y_train[i]) # Store corresponding labels
plt.show()

# Cluster the descriptors using MiniBatchKMeans to create a visual vocabulary
des_array = np.vstack(des_list) if des_list else np.empty((0, 128))
k = 10 * len(np.unique(y_train)) # Determine the number of clusters
kmeans = MiniBatchKMeans(n_clusters=k, batch_size=max(des_array.shape[0] // 4,
1), n_init=4)
kmeans.fit(des_array)

# Create histograms of visual words for each training image
x_train_hist = [np.bincount(kmeans.predict(des), minlength=k) / len(des) if des
is not None else np.zeros(k) for des in des_list]

# Standardize features before training the SVM
scaler = StandardScaler()
x_train_scaled = scaler.fit_transform(x_train_hist)

# Configure and run grid search to find the best SVM parameters
param_grid = {'C': [0.1, 1, 10, 100], 'gamma': ['scale', 'auto'], 'kernel':
['linear', 'rbf', 'poly']}
grid_search = GridSearchCV(SVC(class_weight='balanced'), param_grid, cv=5,
verbose=3)

```

```

grid_search.fit(x_train_scaled, y_train_list)

# Compute and print the training accuracy
y_train_pred = grid_search.best_estimator_.predict(x_train_scaled)
print("Training Accuracy:", accuracy_score(y_train_list, y_train_pred))

# Prepare and predict on the validation set
x_val_hist = []
for img in x_val:
    img = img_as_ubyte(img) # Convert the image for SIFT processing
    _, val_des = sift.detectAndCompute(img, None) # Compute descriptors for
validation images
    val_hist = np.bincount(kmeans.predict(val_des), minlength=k) / len(val_des)
if val_des is not None else np.zeros(k)
    x_val_hist.append(val_hist)
x_val_scaled = scaler.transform(x_val_hist) # Scale the histograms

# Predict and evaluate the validation set
y_val_pred = grid_search.best_estimator_.predict(x_val_scaled)
print("Validation Accuracy:", accuracy_score(y_val, y_val_pred))
print("Confusion Matrix (Validation):\n", confusion_matrix(y_val, y_val_pred))
print("Classification Report (Validation):\n", classification_report(y_val,
y_val_pred))

# Plot the confusion matrix for the validation set
plt.figure(figsize=(6, 6))
sns.heatmap(confusion_matrix(y_val, y_val_pred), annot=True, fmt='g',
cmap='viridis', xticklabels=np.unique(y_val), yticklabels=np.unique(y_val))
plt.xlabel('Predicted labels')
plt.ylabel('True labels')
plt.title('Confusion Matrix for Validation Set')
plt.show()

# Save the trained models and scaler to disk for later use
joblib.dump(grid_search.best_estimator_, 'best_model_sift_svm.pkl') # Save the
best SVM model
print("SIFT_svm model saved successfully!")
joblib.dump(kmeans, 'kmeans_sift_svm.pkl') # Save the k-means model
print("KMeans model saved successfully!")
joblib.dump(scaler, 'scaler_sift_svm.pkl') # Save the scaler
print("Scaler saved successfully!")

# Import necessary libraries
import cv2
import numpy as np
from skimage import img_as_ubyte # Convert images to uint8 format
import matplotlib.pyplot as plt
import seaborn as sns # For plotting the confusion matrix
from sklearn.metrics import accuracy_score, confusion_matrix,
classification_report # For model evaluation
import joblib # For loading the trained model

# Load the trained SVM model
model = joblib.load('best_model_sift_svm.pkl') # Ensure this matches the file

```


name used when saving the model

```
# Initialize SIFT detector
sift = cv2.SIFT_create()
des_list_test = [] # List to store descriptors of test images

# Extract SIFT descriptors for each test image
for img in x_test:
    img_uint8 = img_as_ubyte(img) # Convert images to uint8 format
    kp, des = sift.detectAndCompute(img_uint8, None) # Detect keypoints and
compute descriptors
    if des is not None:
        des_list_test.append(des.astype(np.float32)) # Convert descriptors to
float32 to match the training data type
    else:
        des_list_test.append(np.zeros((1, 128), dtype=np.float32)) # Handle
images with no keypoints/descriptors

# Create histograms of visual words for each test image
x_test_hist = np.array([np.bincount(kmeans.predict(des),
minlength=kmeans.n_clusters) / len(des) if des.size > 0 else
np.zeros(kmeans.n_clusters) for des in des_list_test])

# Scale features using the same scaler as training data
x_test_scaled = scaler.transform(x_test_hist)

# Predict labels for test data
y_test_pred = model.predict(x_test_scaled)

# Model evaluation on the test set
print("Test Accuracy:", accuracy_score(y_test, y_test_pred))
print("Confusion Matrix (Test):\n", confusion_matrix(y_test, y_test_pred))
print("Classification Report (Test):\n", classification_report(y_test,
y_test_pred))

# Plot the confusion matrix for visualization
plt.figure(figsize=(8, 8))
sns.heatmap(confusion_matrix(y_test, y_test_pred), annot=True, fmt='g',
cmap='viridis', xticklabels=np.unique(y_test), yticklabels=np.unique(y_test))
plt.xlabel('Predicted labels')
plt.ylabel('True labels')
plt.title('Confusion Matrix for Test Set')
plt.show()
```

CNN-

```
import numpy as np
import matplotlib.pyplot as plt
import os
from PIL import Image
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.applications import MobileNetV2
from tensorflow.keras.layers import GlobalAveragePooling2D, Dense
from tensorflow.keras.models import Model
```

```

from tensorflow.keras.optimizers import Adam
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, confusion_matrix,
classification_report
from tensorflow.keras.callbacks import ModelCheckpoint

# Directories for images and labels
images_dir = '/content/drive/MyDrive/Colab Notebooks/Computer
Vision/CW_Folder_PG/CW_Dataset/train/images'
labels_dir = '/content/drive/MyDrive/Colab Notebooks/Computer
Vision/CW_Folder_PG/CW_Dataset/train/labels'

# Arrays to store images and labels
images = []
labels = []

# Load and preprocess images and labels
for filename in os.listdir(images_dir):
    if filename.lower().endswith(('.jpg', '.jpeg', '.png')):
        # Load and preprocess each image
        img_path = os.path.join(images_dir, filename)
        img = Image.open(img_path).convert('RGB')
        img = img.resize((128, 128))
        img_array = np.array(img) / 127.5 - 1
        images.append(img_array)

        # Load corresponding label
        label_filename = os.path.splitext(filename)[0] + '.txt'
        label_path = os.path.join(labels_dir, label_filename)
        with open(label_path, 'r') as file:
            label = int(file.read().strip())
            labels.append(label)

# Convert list of images and labels to numpy arrays
images = np.array(images)
labels = np.array(labels)

# Split dataset into training and validation sets
x_train, x_val, y_train, y_val = train_test_split(images, labels, test_size=0.2,
random_state=42)

# Data augmentation configuration
data_augmentation = ImageDataGenerator(
    rotation_range=15,          # Degree range for random rotations
    width_shift_range=0.2,      # Fraction of total width to shift images
    height_shift_range=0.2,     # Fraction of total height to shift images
    shear_range=0.2,            # Shear intensity (shear angle in radians)
    zoom_range=0.2,             # Range for random zoom
    horizontal_flip=True,       # Randomly flip inputs horizontally
    fill_mode='nearest'        # Points outside the boundaries of the input are
    filled
)

```

```

# MobileNetV2 as the base model for transfer learning
base_model = MobileNetV2(input_shape=(128, 128, 3), include_top=False,
weights='imagenet', pooling='avg')
base_model.trainable = False

# Adding custom layers on top of the base model
x = base_model.output
predictions = Dense(len(np.unique(labels)), activation='softmax')(x)
model = Model(inputs=base_model.input, outputs=predictions)

# Compile the model
model.compile(optimizer=Adam(learning_rate=0.001),
loss='sparse_categorical_crossentropy', metrics=['accuracy'])

# Model checkpoint to save the best model
checkpoint = ModelCheckpoint(
    'best_CNNmodel.h5', # Path where the model is saved
    monitor='val_accuracy', # Metric name to monitor
    save_best_only=True, # Save only the best model
    verbose=1 # Verbosity level
)

# Use data augmentation generator to train the model
train_datagen = data_augmentation.flow(x_train, y_train, batch_size=32)

# Train the model using the checkpoint
history = model.fit(
    train_datagen,
    steps_per_epoch=len(x_train) / 32,
    epochs=10,
    validation_data=(x_val, y_val),
    callbacks=[checkpoint]
)

# Load the best model
model.load_weights('best_CNNmodel.h5')

# Evaluation on training data
y_train_pred = np.argmax(model.predict(x_train), axis=1)
train_accuracy = accuracy_score(y_train, y_train_pred)
print("Training Accuracy:", train_accuracy)

# Evaluation on validation data
y_val_pred = np.argmax(model.predict(x_val), axis=1)
val_accuracy = accuracy_score(y_val, y_val_pred)
print("Validation Accuracy:", val_accuracy)

# Print confusion matrix and classification report
conf_matrix = confusion_matrix(y_val, y_val_pred)
print("Confusion Matrix:\n", conf_matrix)
class_report = classification_report(y_val, y_val_pred)
print("Classification Report:\n", class_report)

```

```

# Plot the confusion matrix
plt.figure(figsize=(6, 6))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='viridis',
xticklabels=np.unique(y_val), yticklabels=np.unique(y_val))
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.title('Confusion Matrix for the Validation Set')
plt.show()

# Plot training and validation loss and accuracy
plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 1)
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Training and Validation Loss')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Training and Validation Accuracy')
plt.legend()

plt.show()

# Display the model's architecture
model.summary()

import numpy as np # Import NumPy library for numerical computations
import matplotlib.pyplot as plt # Import matplotlib library for data
visualization
import os # Import os module for interacting with the operating system
from PIL import Image # Import Image module from PIL library for image
processing
from tensorflow.keras.preprocessing.image import ImageDataGenerator # Import
ImageDataGenerator for data augmentation
from tensorflow.keras.models import load_model # Import load_model to load
pre-trained models
from sklearn.metrics import accuracy_score, confusion_matrix,
classification_report # Import metrics for model evaluation
import seaborn as sns # Import seaborn for enhanced data visualization

# Path to test images and labels
test_images_dir = '/content/drive/MyDrive/Colab Notebooks/Computer
Vision/CW_Folder_PG/CW_Dataset/test/images'
test_labels_dir = '/content/drive/MyDrive/Colab Notebooks/Computer
Vision/CW_Folder_PG/CW_Dataset/test/labels'

# Load and preprocess test images and labels
test_images = [] # Initialize an empty list to store test images
test_labels = [] # Initialize an empty list to store corresponding test labels

# Loop through test image directory
for filename in os.listdir(test_images_dir):

```

```

    if filename.lower().endswith(('.jpg', '.jpeg', '.png')):
        # Load and preprocess each image
        img_path = os.path.join(test_images_dir, filename) # Construct full
image path
        img = Image.open(img_path).convert('RGB') # Open and convert image to
RGB format
        img = img.resize((128, 128)) # Resize image to (128, 128)
        img_array = np.array(img) / 127.5 - 1 # Convert image to array and
normalize
        test_images.append(img_array) # Append preprocessed image to list

        # Load corresponding label
        label_filename = os.path.splitext(filename)[0] + '.txt' # Generate
label file name
        label_path = os.path.join(test_labels_dir, label_filename) # Construct
full label path
        with open(label_path, 'r') as file:
            label = int(file.read().strip()) # Read label from file
        test_labels.append(label) # Append label to list

# Convert lists to numpy arrays
test_images = np.array(test_images) # Convert list of images to numpy array
test_labels = np.array(test_labels) # Convert list of labels to numpy array

# Load the best model
model = load_model('best_CNNmodel.h5') # Load pre-trained CNN model

# Make predictions on the test set
test_predictions = model.predict(test_images) # Make predictions on test images
test_pred_classes = np.argmax(test_predictions, axis=1) # Get predicted classes

# Evaluate the model on the test set
test_accuracy = accuracy_score(test_labels, test_pred_classes) # Calculate test
accuracy
print("Test Accuracy:", test_accuracy) # Print test accuracy

# Generate the confusion matrix and classification report
test_conf_matrix = confusion_matrix(test_labels, test_pred_classes) # Generate
confusion matrix
test_class_report = classification_report(test_labels, test_pred_classes) #
Generate classification report

# Display the confusion matrix
plt.figure(figsize=(6, 6)) # Set figure size
sns.heatmap(test_conf_matrix, annot=True, fmt='d', cmap='viridis',
xticklabels=np.unique(test_labels), yticklabels=np.unique(test_labels))
plt.xlabel('Predicted Labels') # Set x-axis label
plt.ylabel('True Labels') # Set y-axis label
plt.title('Confusion Matrix for the Test Set') # Set title
plt.show() # Show plot

# Print the classification report
print("Classification Report:\n", test_class_report) # Print classification
report

```

```

import matplotlib.pyplot as plt
from itertools import cycle
from sklearn.metrics import roc_curve, auc
from sklearn.preprocessing import label_binarize
import numpy as np

# Binarize the labels for multi-class ROC analysis
y_bin = label_binarize(test_labels, classes=np.unique(test_labels))
n_classes = y_bin.shape[1]

# Compute ROC curve and ROC area for each class
fpr = dict()
tpr = dict()
roc_auc = dict()
for i in range(n_classes):
    fpr[i], tpr[i], _ = roc_curve(y_bin[:, i], test_probabilities[:, i])
    roc_auc[i] = auc(fpr[i], tpr[i])

# Plot all ROC curves
plt.figure(figsize=(10, 8), facecolor='black') # Set the background to black
ax = plt.gca()
ax.set_facecolor('black') # Set the plot background to black
colors = cycle(['aqua', 'darkorange', 'cornflowerblue', 'lime', 'magenta',
'gold', 'teal', 'white'])
for i, color in zip(range(n_classes), colors):
    plt.plot(fpr[i], tpr[i], color=color, lw=2,
             label='ROC curve of class {0} (area = {1:0.2f})'.format(i,
roc_auc[i]))

plt.plot([0, 1], [0, 1], 'k--', lw=2)
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate', fontsize=16, color='white')
plt.ylabel('True Positive Rate', fontsize=16, color='white')
plt.title('Extended ROC Curves for Each Class', fontsize=18, color='white')
plt.legend(loc="lower right")
plt.grid(True, color='gray')
plt.xticks(np.arange(0.0, 1.1, step=0.1), color='white')
plt.yticks(np.arange(0.0, 1.1, step=0.1), color='white')
plt.tick_params(axis='both', colors='white')
ax.spines['top'].set_color('white')
ax.spines['bottom'].set_color('white')
ax.spines['right'].set_color('white')
ax.spines['left'].set_color('white')
plt.gca().set_aspect('equal', adjustable='box')
plt.show()

```